

Reinforcement Learning (& Decision-Making) with hBayesDM



*Woo-Young (Young) Ahn
Department of Psychology
Seoul National University
ccs-lab.github.io*

General Outline

*Part I: Introduction to computational modeling and
hBayesDM (hierarchical Bayesian modeling of
Decision-Making tasks)*



*Part II: Hands-on tutorials on Bayesian modeling /
hBayesDM*



Nate Haines

github.com/CCS-Lab/cpc2018

CCS-Lab / cpc2018

Unwatch 1 Star 0 Fork 0

Code Issues 0 Pull requests 0 Projects 0 Wiki Insights Settings

hBayesDM workshop at Computational Psychiatry Course (CPC) 2018 Edit

cpc2018 hbayesdm workshop Manage topics

3 commits 1 branch 0 releases 1 contributor GPL-3.0

Branch: master New pull request Create new file Upload files Find file Clone or download ▾

youngahn upload R codes

Data upload R codes

R_Codes/tutorial upload R codes

LICENSE Initial commit

README.md Update README.md

Clone with HTTPS Use SSH
Use Git or checkout with SVN using the web URL.
<https://github.com/CCS-Lab/cpc2018.git>

Open in Desktop Download ZIP

README.md

git clone <https://github.com/CCS-Lab/cpc2018.git>

Learning objectives

Participants will...

- *Know popular and cutting-edge methods for fitting computational models*
- *Review the concept of Bayesian data analysis*
- *Evaluate outputs from Bayesian data analysis*
- *Use hBayesDM package to model reinforcement learning and decision-making tasks*

Outline for Part I

- *What is computational modeling (a.k.a. cognitive modeling)?*
- *Why/how do we lower the barrier to computational modeling?*
 - *Brief introduction to hBayesDM*
- *How to fit a computational model?*
 - *Maximum likelihood estimation (MLE)*
 - *Bayesian analysis & MCMC sampling*
 - *Hierarchical Bayesian analysis*
 - *Tools for Bayesian data analysis*
- *Things to know when performing MCMC sampling*

Software installation (1/2)

- *Install the latest R (+3.4):* <https://cloud.r-project.org/>
- *Install RStudio:*
<https://www.rstudio.com/products/rstudio/download2/>

Windows users

- *Install Rtools:* <https://cran.r-project.org/bin/windows/Rtools/Rtools34.exe>
- **“Carefully” follow the instructions:** <https://github.com/stan-dev/rstan/wiki/Install-Rtools-for-Windows>

Mac users

- *Install Xcode* (<https://developer.apple.com/xcode/>)

Software installation (2/2)

hBayesDM tutorial: <http://rpubs.com/CCSL/hBayesDM>

“How to install hBayesDM” (Google “hBayesDM rpubs”)

Windows users

- *Install hBayesDM from CRAN. Takes ~10 seconds.*

```
install.packages("hBayesDM", dependencies=TRUE)
```

Mac/Linux users

- *Install hBayesDM from GitHub. Takes ~5-10 minutes.*

```
# install 'devtools' if required
if (!require(devtools)) install.packages("devtools")
devtools::install_github("CCS-Lab/hBayesDM")
```

Outline

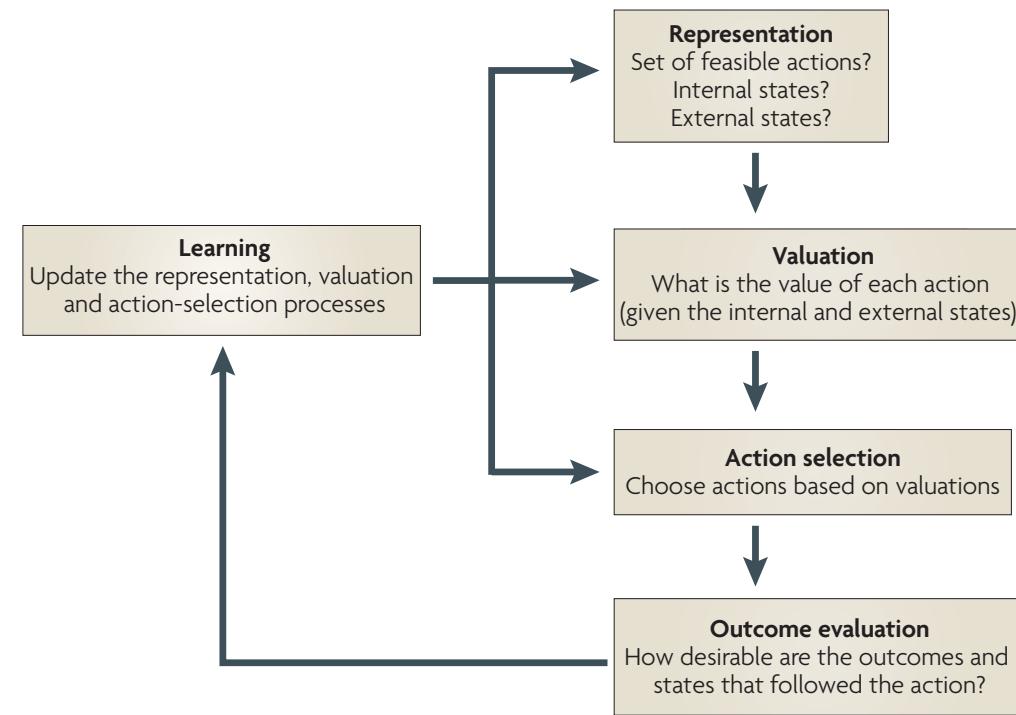
- *Understand the concepts of computational modeling*
- *Know popular and cutting-edge methods for fitting computational models / Understand the concept of Bayesian data analysis*
- *Things to know when performing MCMC sampling*
- *Use hBayesDM package to model several tasks*

Reinforcement learning and Decision-making (RLDM)

A framework for studying the neurobiology of value-based decision making

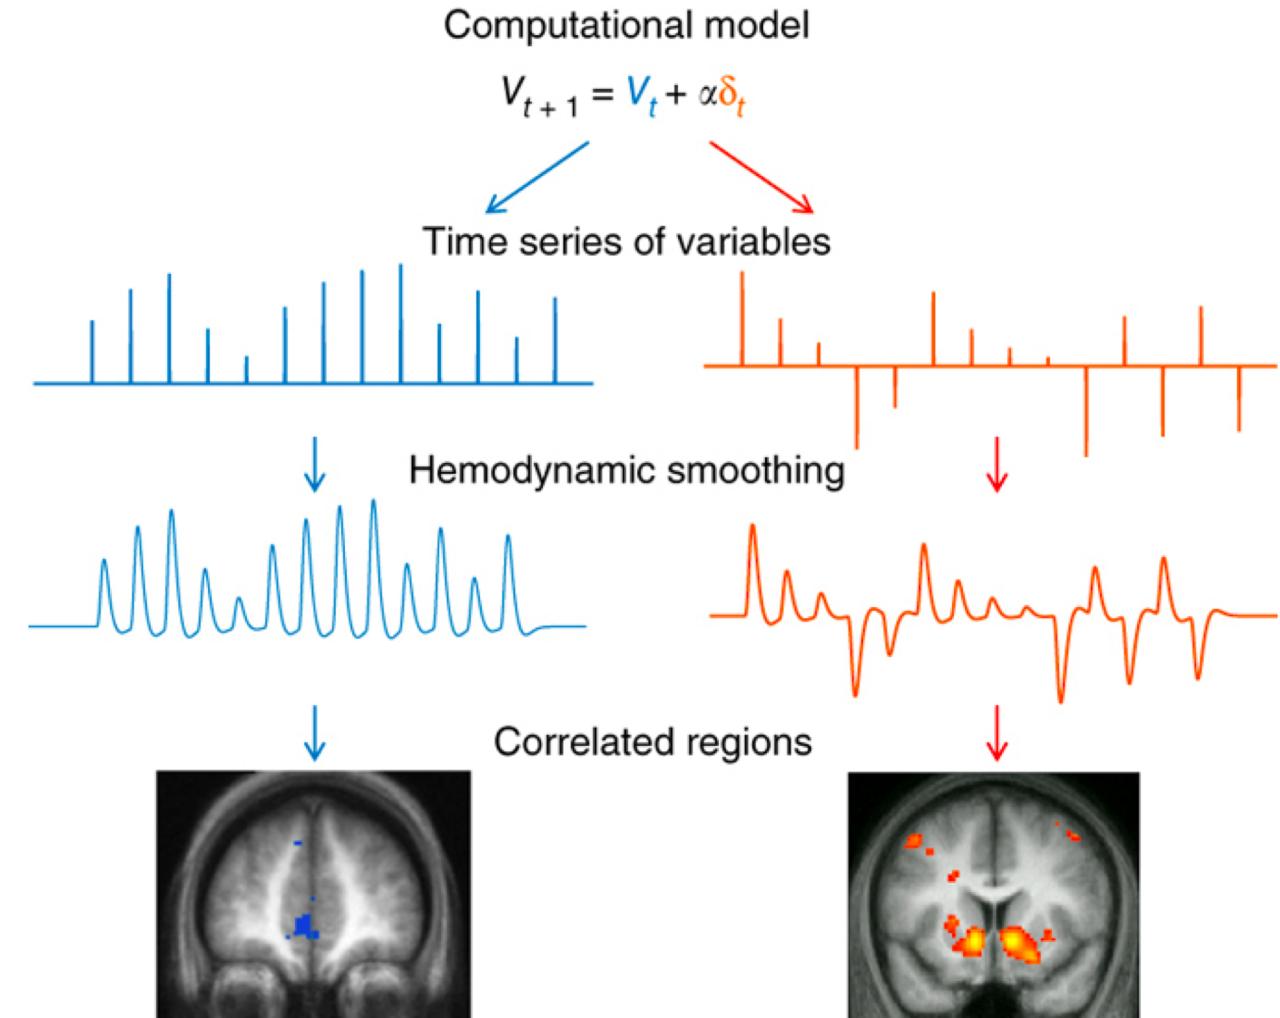
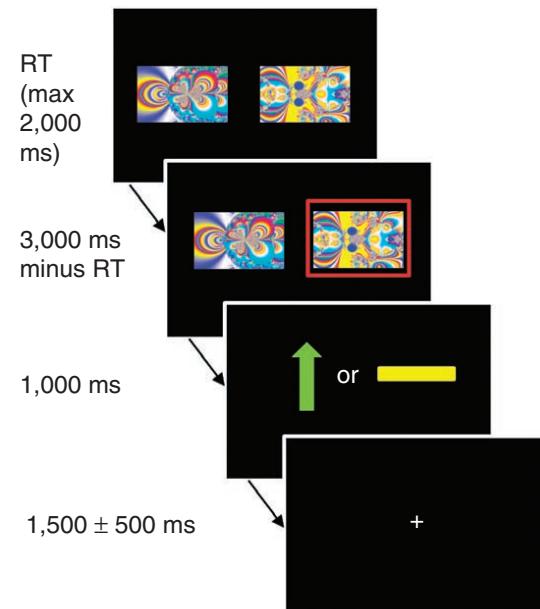
Antonio Rangel*, Colin Camerer* and P. Read Montague†

Rangel et al (2008) *Nature Rev. Neuro*



Model-based fMRI/EEG

e.g., Forstmann & Wagenmakers (2015); O'Doherty et al (2007); Cohen et al (2017)



Special Issue: Cognition in Neuropsychiatric Disorders

Computational psychiatry

P. Read Montague^{1,2}, Raymond J. Dolan², Karl J. Friston² and Peter Dayan³



Computational psychiatry: the brain as a phantastic organ

Karl J Friston, Klaas Enno Stephan, Read Montague, Raymond J Dolan



Available online at www.sciencedirect.com

ScienceDirect

Current Opinion in
Neurobiology

Computational approaches to psychiatry
Klaas Enno Stephan^{1,2,3} and Christoph Mathys³

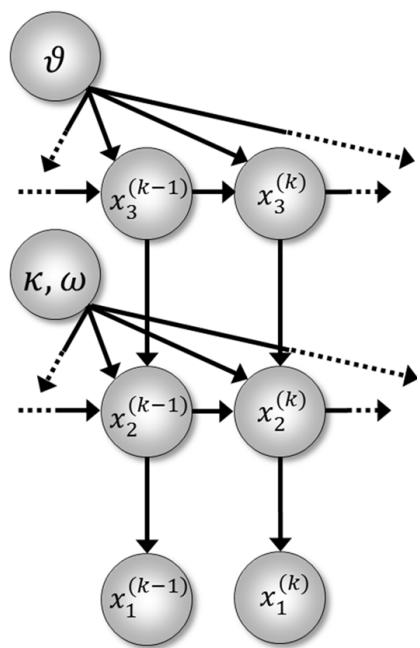
*Computational accounts of
abnormal cognition &
its biological underpinnings*

The screenshot shows the homepage of the journal 'nature' with a red header. The header includes the title 'nature International weekly journal of science' and a navigation bar with links to Home, News & Comment, Research, Careers & Jobs, Current Issue, Archive, Audio & Video, and a search bar. Below the header, a breadcrumb navigation shows the path: Archive > Volume 539 > Issue 7627 > News: Q&A > Article. The main content area features a headline 'NATURE | NEWS: Q&A' and a sub-headline 'US mental-health chief: psychiatry must get serious about mathematics'. There is also a share icon in the bottom right corner.

Limitations & Future directions

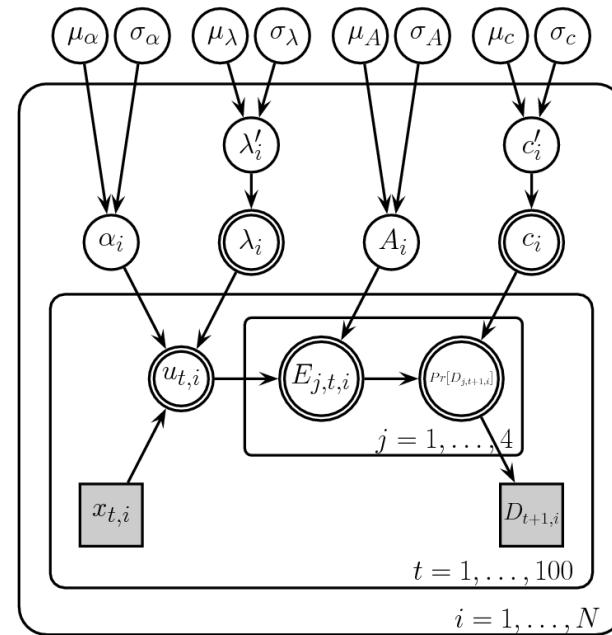
- *Overly simplified “toy” problems*
 - *Violated assumptions* (e.g., *discrete space/action & Markov..*) ,
(Gershman & Daw, 2016, Annu Rev Psych)
 - *Predict real-life DM?* (*Mobb et al., 2018, Nat Rev Neuro*)
- *One-shot learning with sparse data*
 - *Episodic memory (hippocampus)* (*Gabrieli, 1998; Eichenbaum et al., 1999*)
- *Adaptive Design Optimization (ADO)*
- *Modeling of even toy problems is hard for many people*

I like the idea of modeling



Mathys et al (2011) Frontiers

But...

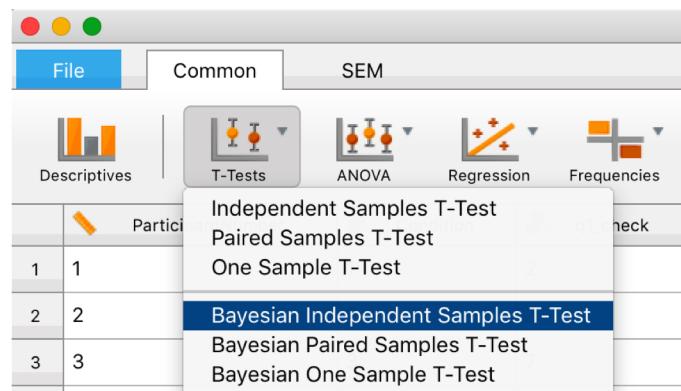


Ahn et al (2011) JNPE

Can we make it easy to do computational modeling?

Q) As easy as doing a *T-test*?

JASP, SPSS



R

```
Console ~/Desktop/ ↗  
> t.test(group1, group2)
```

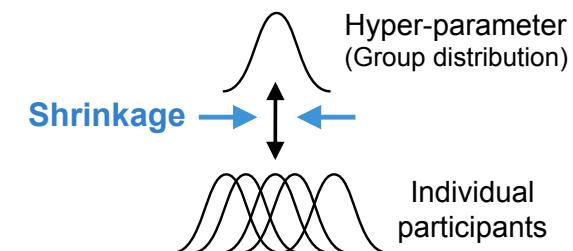
Several packages exist, but ...

Matze et al (2013) Frontiers; Wabersich & Vandekerckhove (2014); Wiecki et al (2013) Frontiers;
Daunizeau et al (2014) PLoS Comp Biol

hBayesDM (hierarchical Bayesian modeling of Decision-Making tasks) Package

- *Models for “many” tasks/paradigms (next slide)*
- *Single-line of coding in R*
 - *Model fitting, visualization, model comparisons*
- *Based on the advanced Bayesian software, Stan (<https://mc-stan.org>).*
- *Hierarchical Bayesian modeling*
- *All codes are publicly available*

<https://github.com/CCS-Lab/hBayesDM>



What tasks and models are available? in CPC2017

Ahn & Busemeyer (2016) Curr Opin Behav Sci

- *Choice reaction time* → sequential sampling models
- *Delay Discounting* (e.g., Mazur, 1987)
- *Iowa Gambling* (Bechara et al, 1994)
- *(Orthogonalized) Go/Nogo* (Guitart-Masip et al, 2012)
- *Two-armed Bandit (Experience-based) including Reversal Learning* (e.g., Erev et al, 2010)
- *Four-armed Bandit (Experience-based)* (e.g., Seymour et al, 2012)
- *Two-choice Description-based* (e.g., Sokol-Hessner et al, 2009; Tom et al, 2007)
- *Ultimatum Game* (e.g., Xiang et al, 2013)
- **Two-Step* (Daw et al, 2011)

*Version 0.4.1. or later

*What tasks and models are available *in CPC2018*?*

Ahn & Busemeyer (2016) Curr Opin Behav Sci

- *Balloon Analogue Risk Task (BART)* ([Wallsten et al, 2005](#))
- *Choice under Risk and Ambiguity* ([Levy et al, 2010](#))
- *Choice reaction time* → sequential sampling models
- *Delay Discounting* ([e.g., Mazur, 1987](#))
- *Iowa Gambling* ([Bechara et al, 1994](#))
- *(Orthogonalized) Go/Nogo* ([Guitart-Masip et al, 2012](#))
- *Peer influence task* ([Chung et al, 2015](#))
- *Probabilistic Selection task* ([e.g., Frank et al, 2007](#))
- *Two-armed Bandit (Experience-based) including Reversal Learning* ([e.g., Erev et al, 2010](#))
- *Four-armed Bandit (Experience-based)* ([e.g., Seymour et al, 2012](#))
- *Two-choice Description-based* ([e.g., Sokol-Hessner et al, 2009; Tom et al, 2007](#))
- *Ultimatum Game* ([e.g., Xiang et al, 2013](#))
- *Two Step task* ([Daw et al, 2011](#))
- *Wisconsin Card Sorting task* ([Bishara et al, 2010](#))

To-do list

- More tasks
 - Kalman-filter ([Daw et al., 2006](#))
 - Hierarchical Gaussian Filtering ([Mathys et al., 2011](#))
 - More models for the delay discounting task.
 - More sequential sampling models (e.g., drift diffusion models with different drift rates for multiple conditions).
 - Passive avoidance learning task ([Newman & Kosson, 1986](#); [Newman et al., 1985](#))
 - Description-based tasks with a probability weighting function (e.g., [Erev et al., 2010](#); [Hertwig et al., 2004](#); [Jessup et al., 2008](#))
 - Stop Signal Task (SST)
 - Model-based regressors ([O'Doherty et al., 2007](#)) from more tasks.
- *hBayesDM in Python* → easy to use with Python-based neuroimaging packages (e.g., Nipype)
- GUI interface (RShiny)

~Downloads/hBayesDMh - Shiny

http://127.0.0.1:5610 | Open in Browser |

hBayesDM GUI

You have selected ra _ prospect

plot type: individual

Parameters to Plot: (empty input field)

Task: ra

Model: prospect

For an introduction of models, read the [hBayesDM Models](#)

Use Other Data File

Number of iteration: 2,000 (Slider from 20 to 4,000)

Number of warmup: 1,000 (Slider from 10 to 2,000)

Number of chain: 2 (Slider from 1 to 10)

Number of core: 2 (Slider from 1 to 5)

Number of thin: 1 (Slider from 1 to 5)

modelRegressor

vb

inc_postpred

GO

How can I use it?

Tutorials available at

<http://rpubs.com/CCSL/hBayesDM>

(you can find it by Googling ‘hBayesDM’)

Install it just like other R packages

```
install.packages("hBayesDM", dependencies=TRUE)
```

```
devtools::install_github("CCS-Lab/hBayesDM")
```

Brief step-by-step tutorials

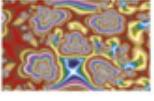
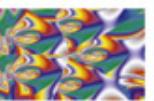
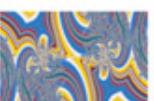
* *Part II for detailed tutorials*

1. *Prepare raw (trial-by-trial) data*
2. *Fit candidate models*
3. *Plot (visualize) and inspect model parameters*
4. *Compare models (if there exist competing models)*

The Orthogonalized Go/Nogo task

- gng_m1
- gng_m2
- gng_m3
- gng_m4

*Guitart-Masip et al, 2012 Neuroimage
Cavanagh et al, 2013 J Neuro*

	punishment	reward
go	go to avoid losing 	goto win 
nogo	nogo to avoid losing 	nogo to win 

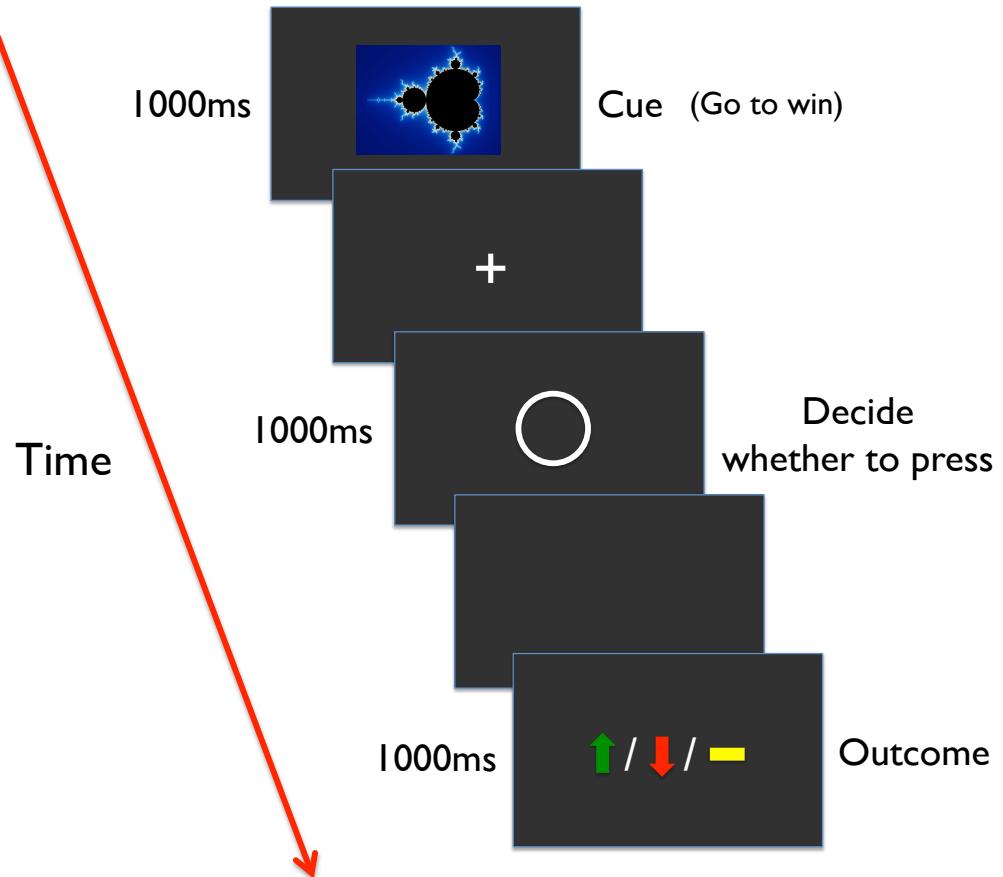
Orthogonalized Go/Nogo task

	Loss	Gain
Go	Go to avoid	Go to win
Nogo	Nogo to avoid	Nogo to win



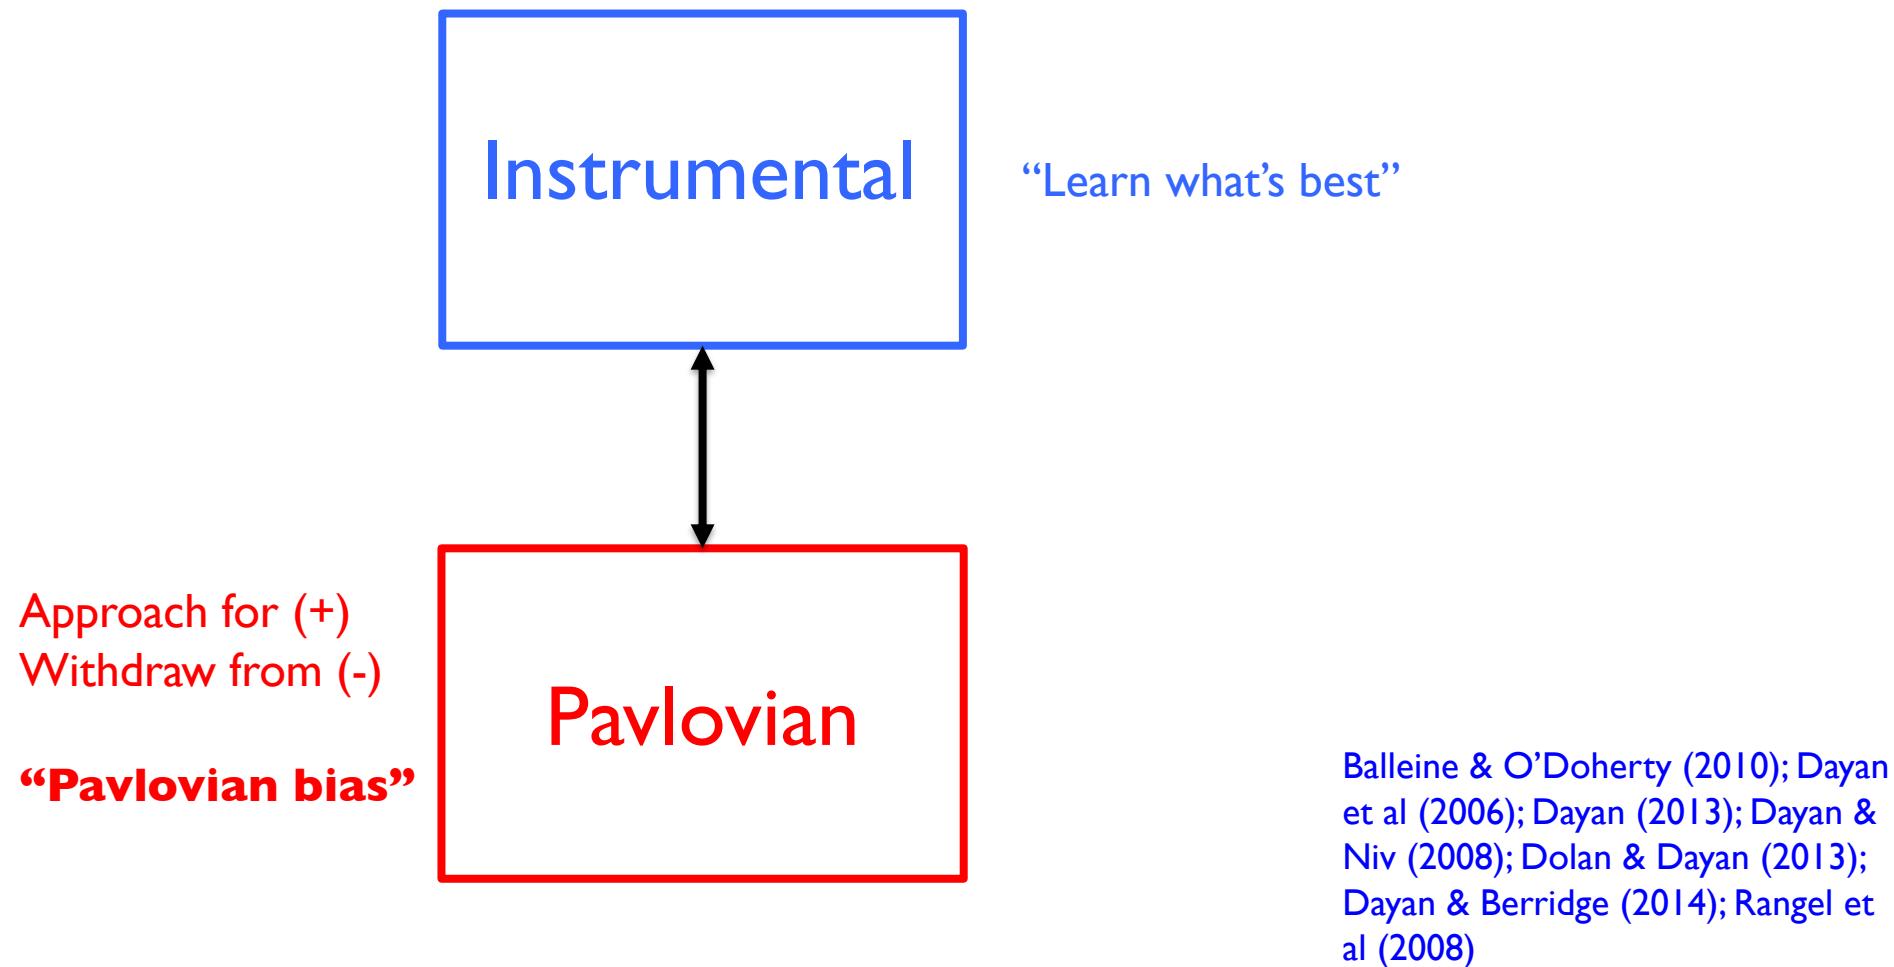
- 4 cues (conditions)

2 actions (Go / Nogo) x
2 valence (Gain / Loss)



Orthogonalized Go/Nogo task

Pavlovian-Instrumental competition



Pavlovian-Instrumental competition

Nogo

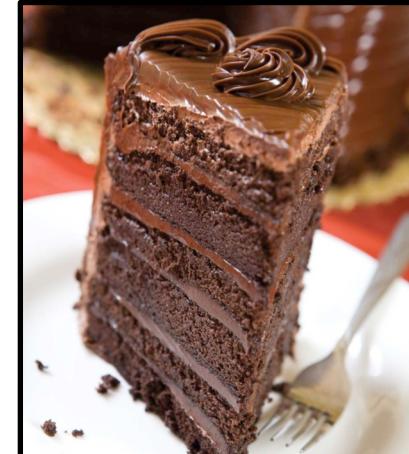
*Instrumental
response*



Young on diet

Go

*Pavlovian
response*



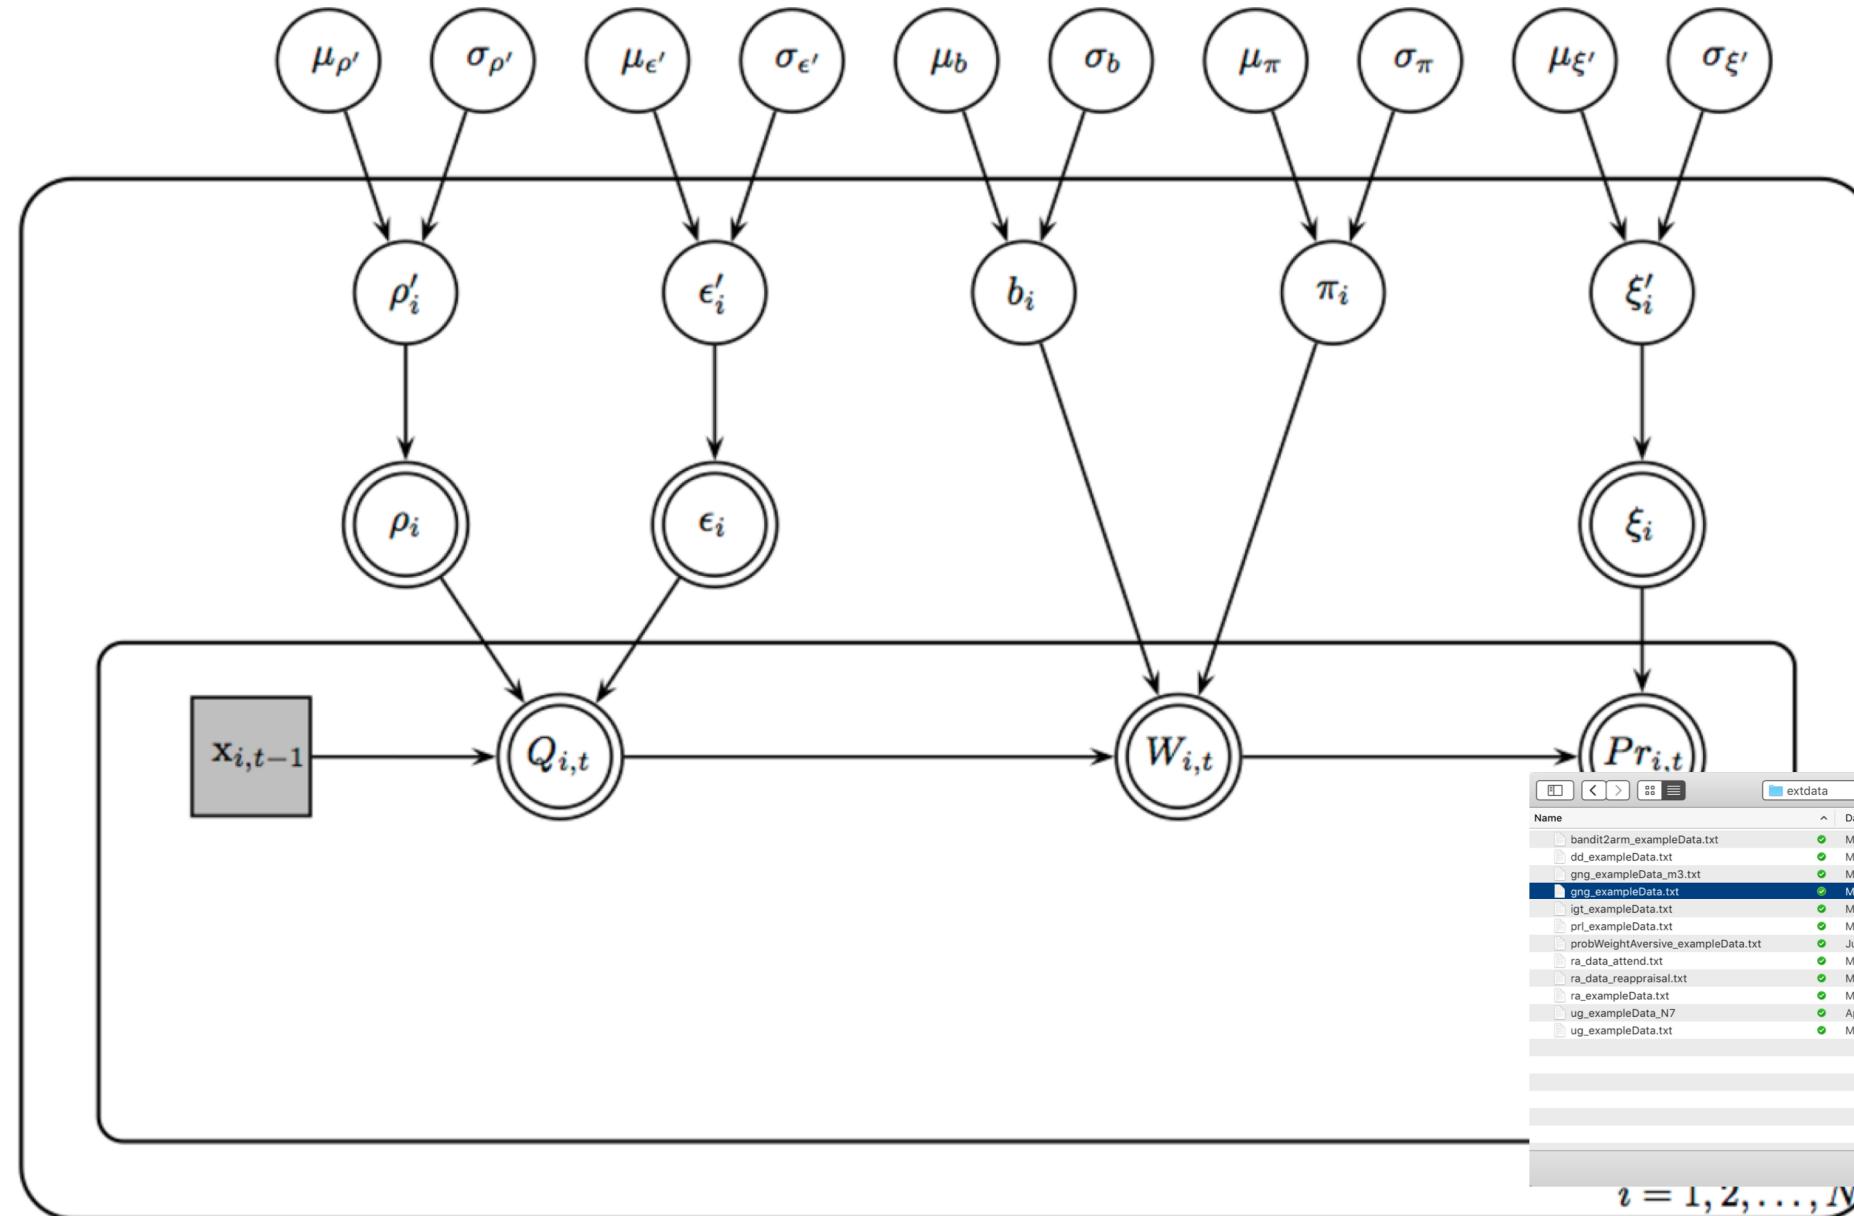
Instrumental goal: To lose weight

1. Prepare raw (trial-by-trial) data as a single text file

A	B	C	D	E	F	G	
1	trialNum	cue	keyPressed	success	congruentOutcome	outcome	subjID
2	1	1	1	1	2	0	1
3	2	2	0	1	1	1	1
4	3	4	0	1	1	0	1
5	4	4	1	0	1	-1	1
6	5	4	0	1	1	0	1
7	6	1	1	1	1	1	1
8	7	3	0	0	1	-1	1
9	8	1	1	1	1	1	1
10	9	3	1	1	1	0	1
11	10	3	0	0	1	-1	1
12	11	4	0	1	1	0	1
13	12	4	0	1	1	0	1
14	13	4	0	1	1	0	1

2. Fit candidate models

output1 =



Name	Date Modified	Size	Kind
bandit2arm_exampleData.txt	May 25, 2016, 12:16 PM	26 KB	Plain Text
dd_exampleData.txt	Mar 13, 2016, 8:15 AM	43 KB	Plain Text
gng_exampleData_m3.txt	Mar 18, 2016, 4:29 PM	38 KB	Plain Text
gng_exampleData.txt	Mar 18, 2016, 4:06 PM	38 KB	Plain Text
igt_exampleData.txt	Mar 16, 2016, 11:16 PM	6 KB	Plain Text
prl_exampleData.txt	May 25, 2016, 12:17 PM	24 KB	Plain Text
probWeightAversive_exampleData.txt	Jul 30, 2016, 1:01 AM	12 KB	Plain Text
ra_data_attend.txt	Mar 15, 2016, 4:25 PM	76 KB	Plain Text
ra_data_reappraisal.txt	Mar 15, 2016, 4:25 PM	77 KB	Plain Text
ra_exampleData.txt	Mar 15, 2016, 4:25 PM	12 KB	Plain Text
ug_exampleData_N7	Apr 13, 2016, 9:03 PM	5 KB	Plain Text
ug_exampleData.txt	May 1, 2016, 2:27 PM	26 KB	Plain Text

$i = 1, 2, \dots, N$

Details:

```
# of chains          = 4
# of cores used     = 4
# of MCMC samples (per chain) = 2000
# of burn-in samples = 1000
# of subjects        = 10
# of (max) trials per subject = 240
```

** Loading a precompiled model **

starting worker pid=75130 on localhost:11950 at 08:25:48.905

starting worker pid=75138 on localhost:11950 at 08:25:49.101

SAMPLING FOR MODEL 'gng_m1' NOW (CHAIN 1).

Chain 1, Iteration: 1 / 2000 [0%] (Warmup)

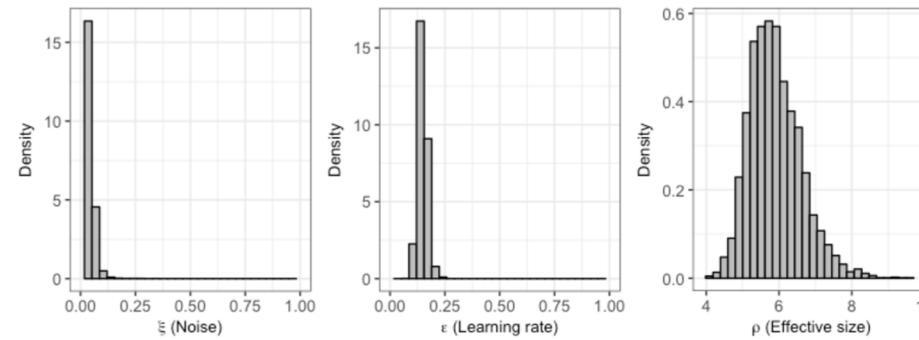
SAMPLING FOR MODEL 'gng_m1' NOW (CHAIN 2).

...

**** Model fitting is complete! ****

3. Plot (visualize) and inspect model parameters

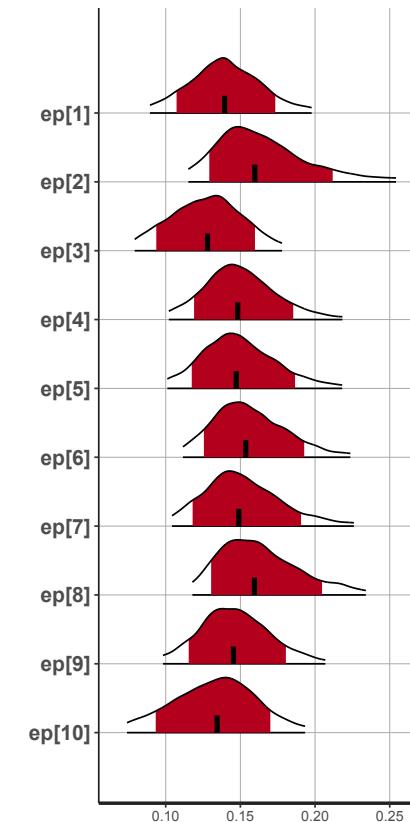
```
plot(output1)
```



```
> output1$allIndPars
```

	xi	ep	rho	subjID
1	0.03688558	0.1397615	5.902901	1
2	0.02934812	0.1653435	6.066120	2
3	0.04467025	0.1268796	5.898099	3
4	0.02103926	0.1499842	6.185020	4
5	0.02620808	0.1498962	6.081908	5
...				

```
plotInd(output1, "ep")
```



4. Bayesian model comparisons

Vehtari et al. (2016)

Leave-One-Out Information Criterion (LOOIC) - default

Widely Applicable Information Criterion (WAIC)

```
> printFit(output1, output2, output3, output4)
```

*Model #4 is the best model
(in terms of LOOIC)*

5. Group comparisons

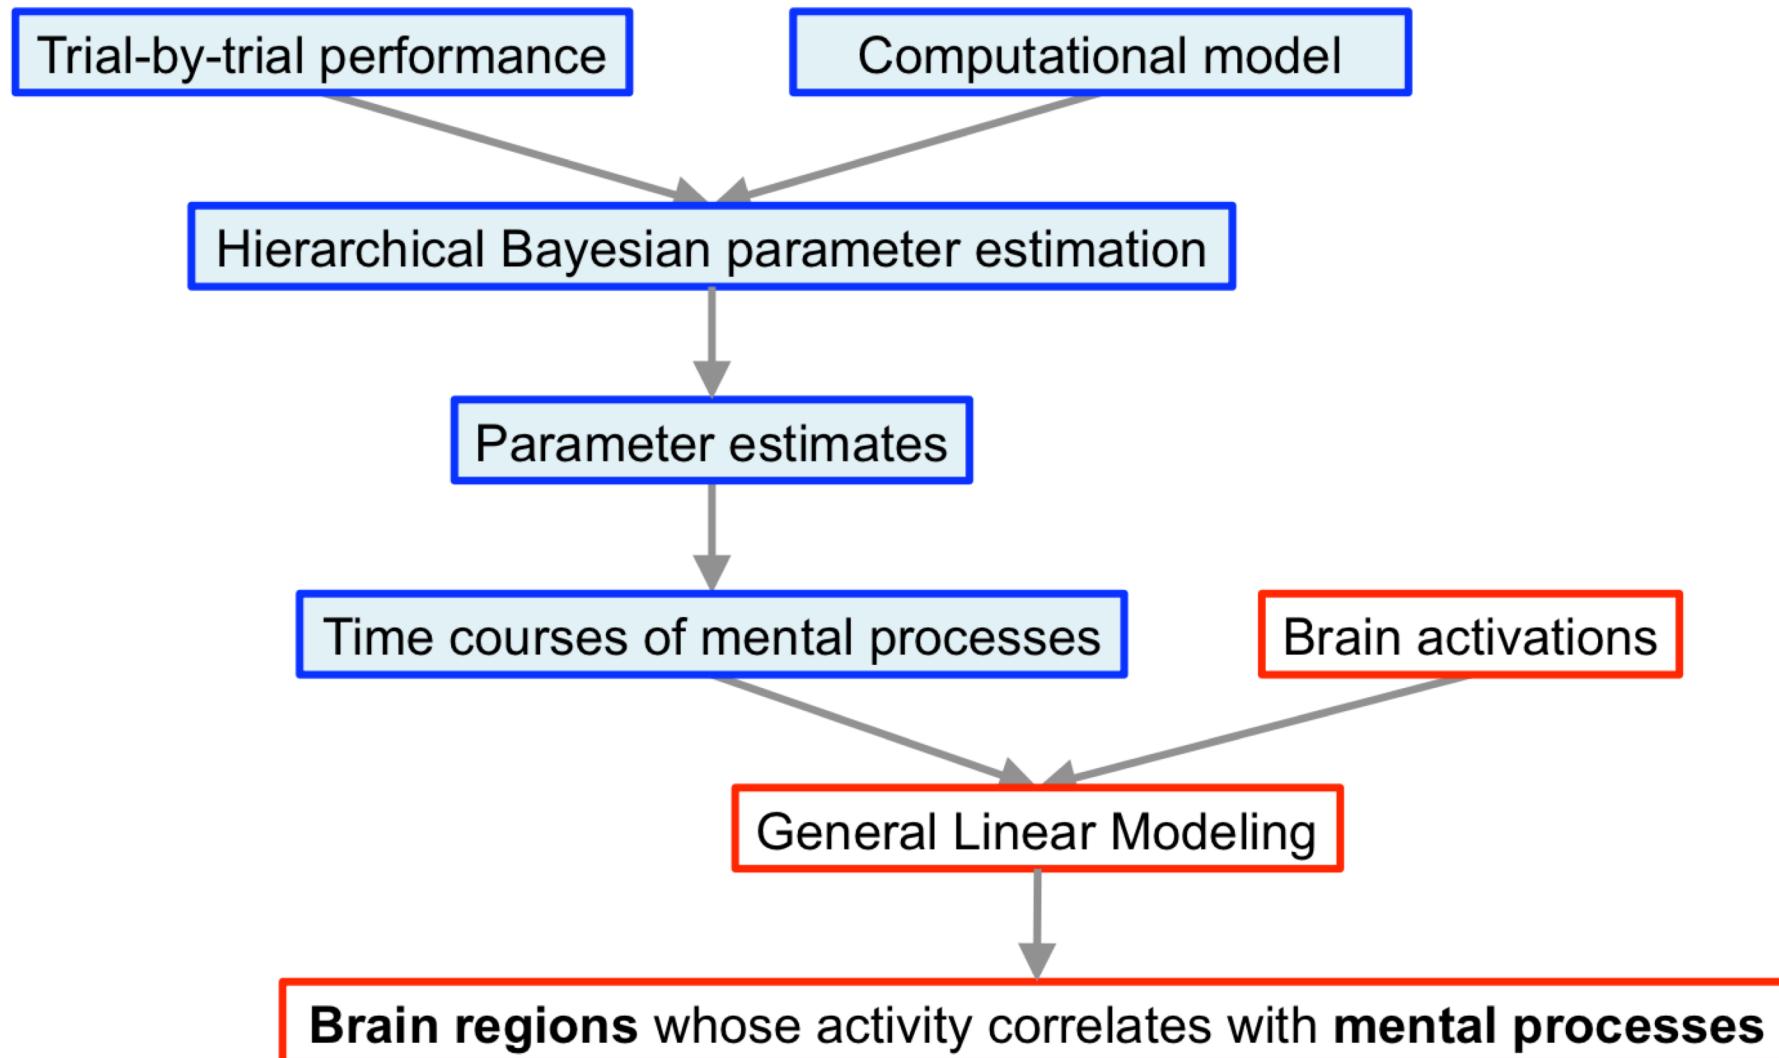
```
data_group1 = "~/Project_folder/gng_data_group1.txt" # data file for group1
data_group2 = "~/Project_folder/gng_data_group2.txt" # data file for group2

output_group1 = gng_m4(data_group1) # fit group1 data with the gng_m4 model
output_group2 = gng_m4(data_group2) # fit group2 data with the gng_m4 model

# After model fitting is complete for both groups,
# evaluate the group difference (e.g., on the 'pi' parameter) by examining the posterior distribution of
# group mean differences.

diffDist = output_group1$parVals$mu_pi - output_group2$parVals$mu_pi # group1 - group2
HDIofMCMC( diffDist ) # Compute the 95% Highest Density Interval (HDI).
plotHDI( diffDist ) # plot the group mean differences
```

6. Model-based fMRI/EEG



Model-based fMRI/EEG

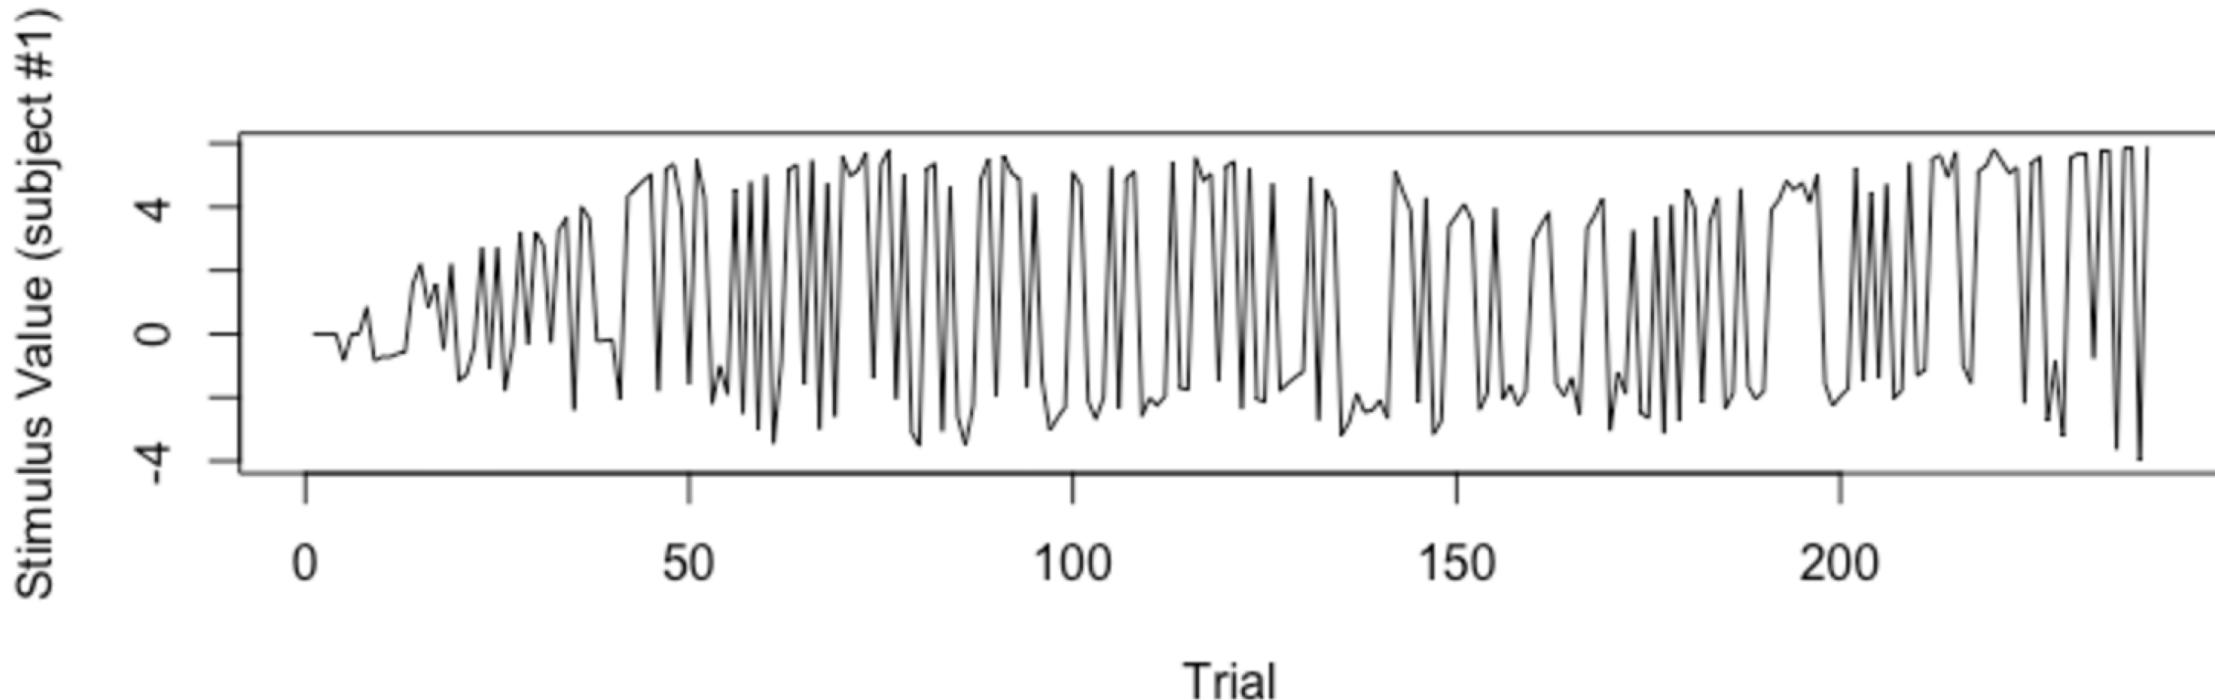
- *Trial-by-trial regressors are readily available*

```
# fit example data with the gng_m3 model
output = gng_m3(data="example", niter=3000, nwarmup=1000, nchain=3, ncore=3, modelRegressor=TRUE)
```

```
# store all subjects' stimulus value (SV) in 'sv_all'  
sv_all = output$modelRegressor$SV  
  
dim(output$modelRegressor$SV) # number of rows=# of subjects (=10), number of columns=# of trials (=240)
```

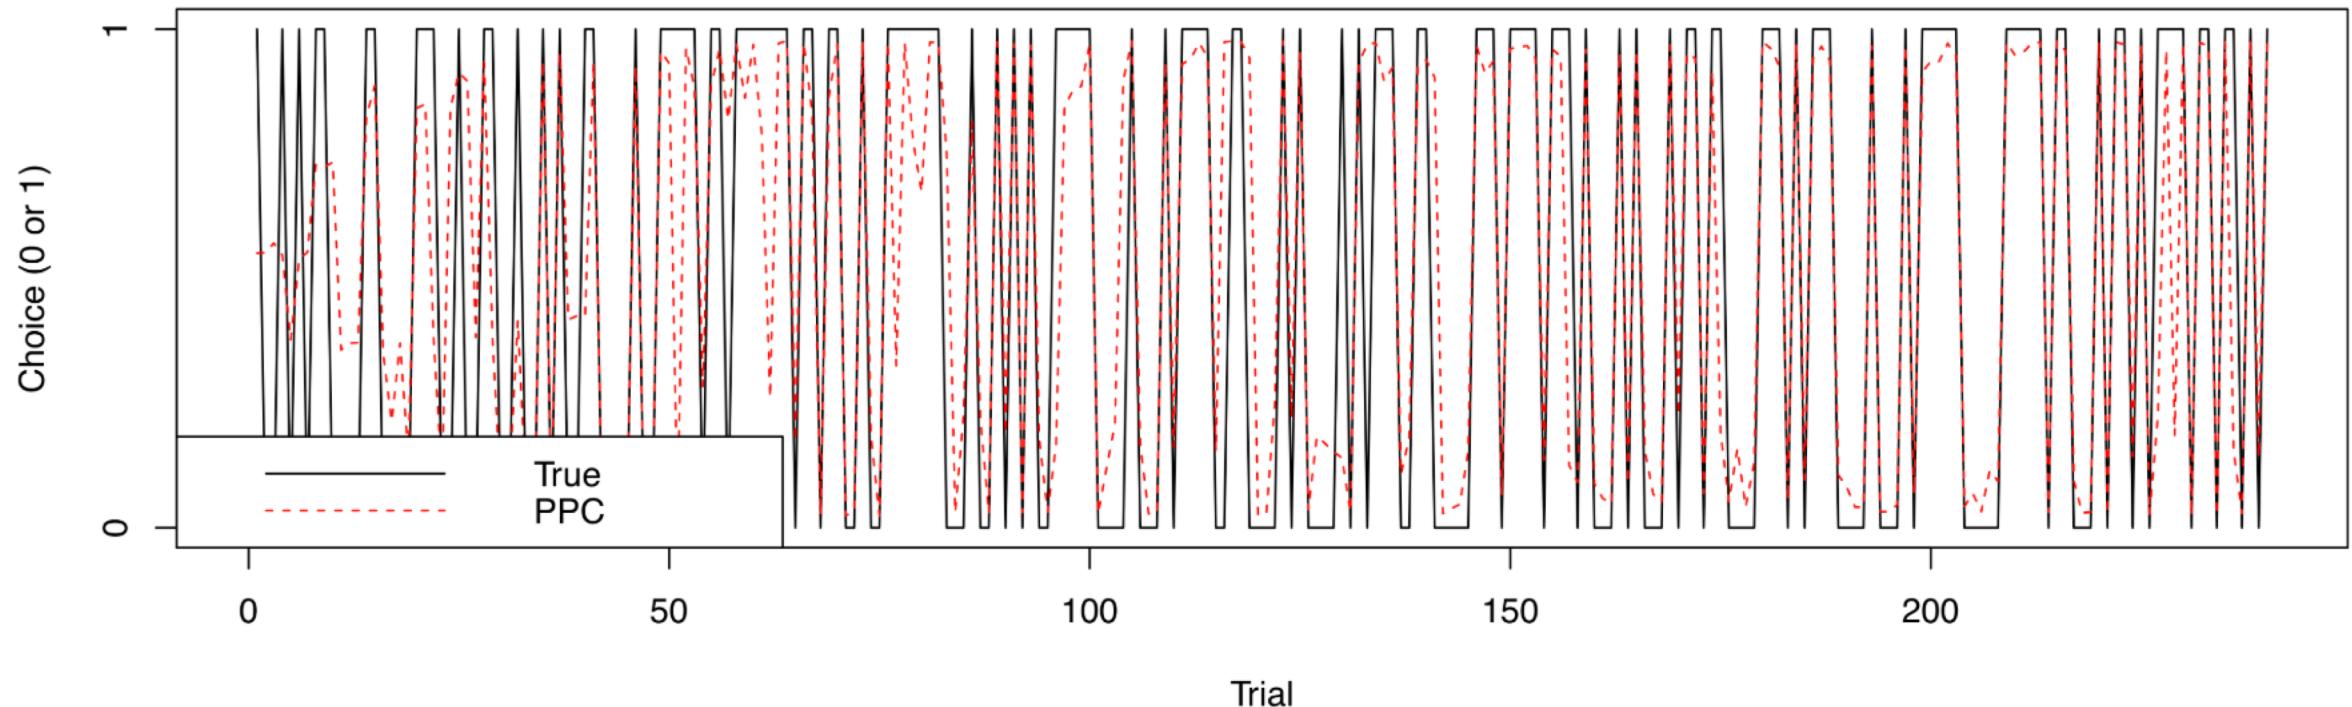
```
## [1] 10 240
```

```
# visualize SV (Subject #1)  
plot(sv_all[1, ], type="l", xlab="Trial", ylab="Stimulus Value (subject #1)")
```



Posterior predictive checks

```
# fit example data with the gng_m3 model and run posterior predictive checks  
x = gng_m3(data="example", niter=2000, nwarmup=1000, nchain=4, ncore=4, inc_postpred = TRUE)
```



Outline

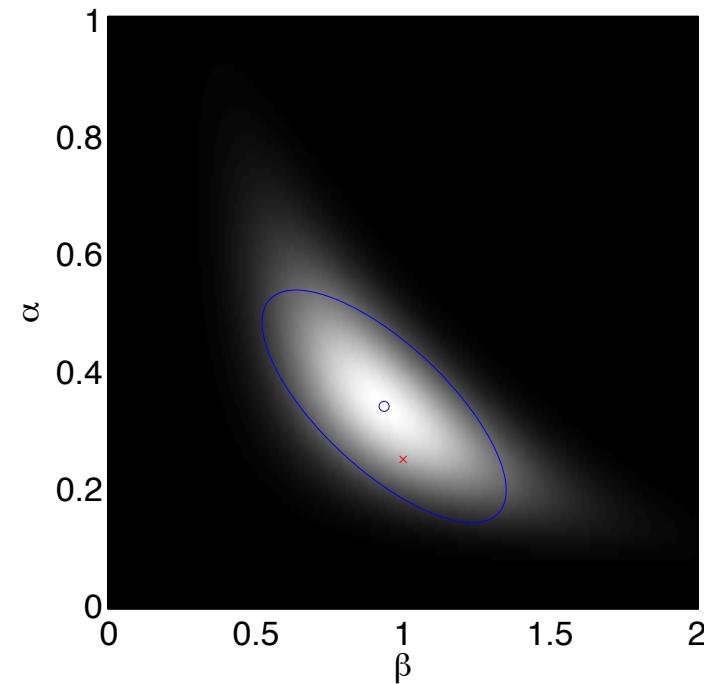
- *Understand the concepts of computational modeling*
- *Know popular and cutting-edge methods for fitting computational models / Understand the concept of Bayesian data analysis*
- *Things to know when performing MCMC sampling*
- *Use hBayesDM package to model several tasks*

Popular methods for model fitting

- *Maximum likelihood estimation (MLE)*
- *Bayesian parameter estimation (or Bayesian data analysis)*
- *Hierarchical Bayesian analysis*

Maximum likelihood estimation

- Ronald A. Fisher (1920s): finds the value of a parameter that makes the observed data “most likely”.
- Likelihood function
 - $L(\text{data} \mid \theta)$: likelihood of getting data given θ
- Analytical solutions often don’t exist
- Nonlinear optimization
 - gradient decent
 - `fminsearch` (Matlab), `optim` (R), etc.
- MLE estimates = point estimates



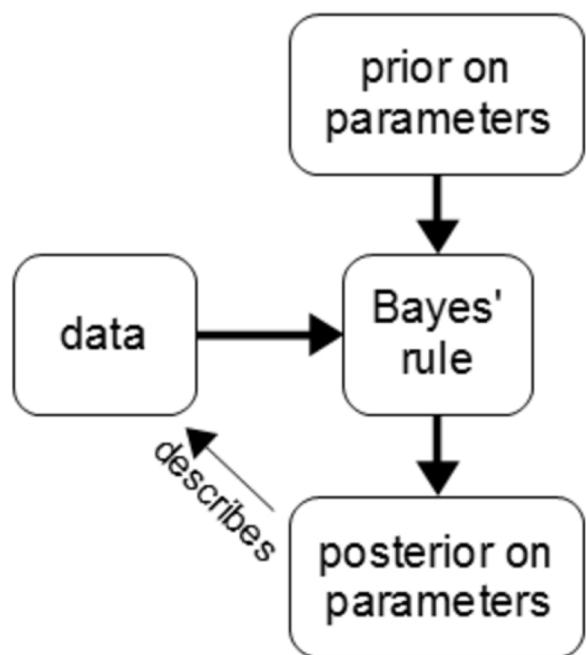
“Likelihood surface” from Daw (2011)

*Brief Introduction to
Bayesian data analysis*

Bayesian data analysis

Bayesian data analysis vs. *Bayesian model of mind*

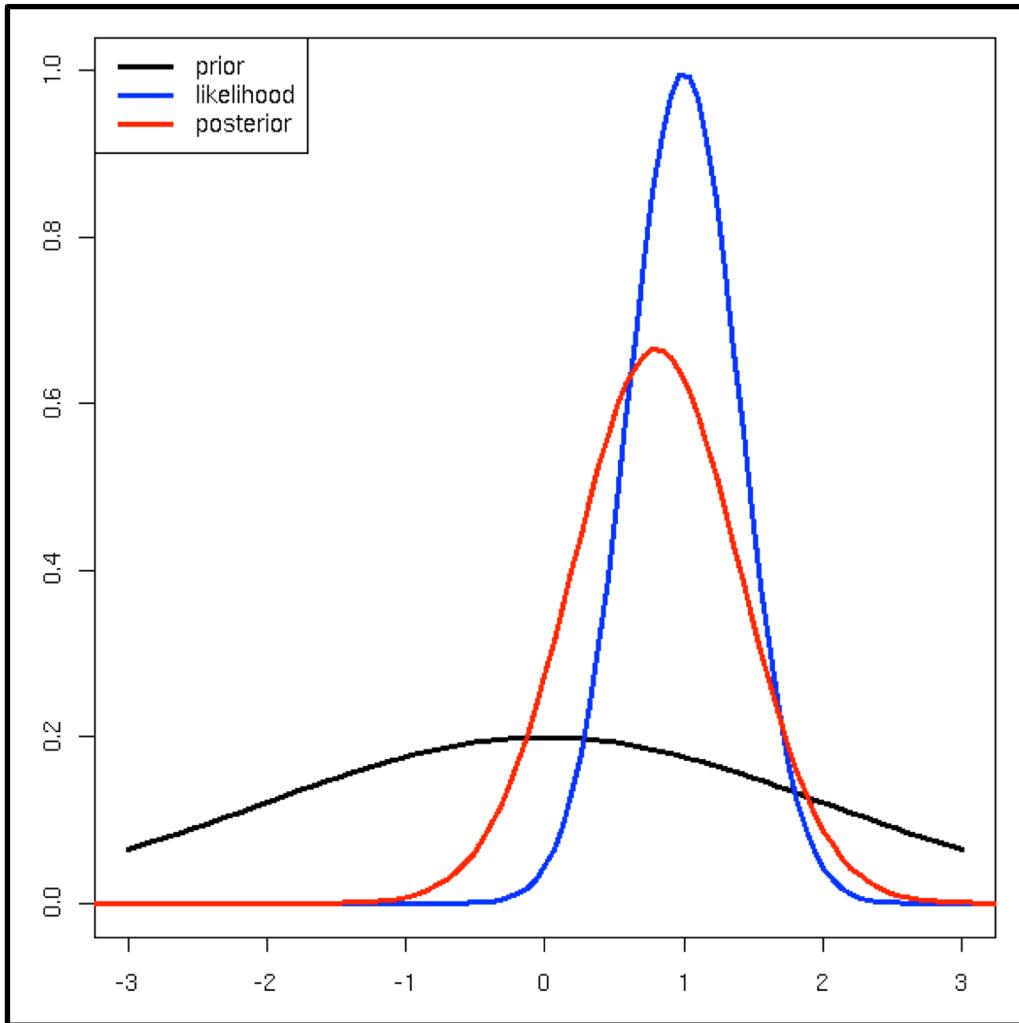
<http://doingbayesiandataanalysis.blogspot.hk/2011/10/bayesian-models-of-mind-psychometric.html>



Θ : Model parameters (e.g., discounting rate, learning rate, etc.)

A diagram illustrating the Bayesian formula. At the top, three boxes are labeled "Posterior", "Likelihood", and "Prior". Below them, the formula is shown: $P(\Theta_i | D_i) = \frac{P(D_i | \Theta_i)P(\Theta_i)}{P(D_i)}$. The term $P(D_i | \Theta_i)P(\Theta_i)$ is enclosed in a red box. The term $P(D_i)$ is also enclosed in a red box. A red arrow labeled "Evidence" points down to the bottom of the formula. A red arrow labeled "Prior" points up to the $P(\Theta_i)$ term. A red arrow labeled "Likelihood" points up to the $P(D_i | \Theta_i)$ term.

$$P(\Theta_i | D_i) = \frac{P(D_i | \Theta_i)P(\Theta_i)}{\int P(D_i | \Theta_{i'})P(\Theta_{i'})d\Theta_{i'}}$$



Bayesian data analysis

Update prior distributions for model parameters into posterior distributions given the data (e.g., trial-by-trial choices and outcomes) using Bayes' rule

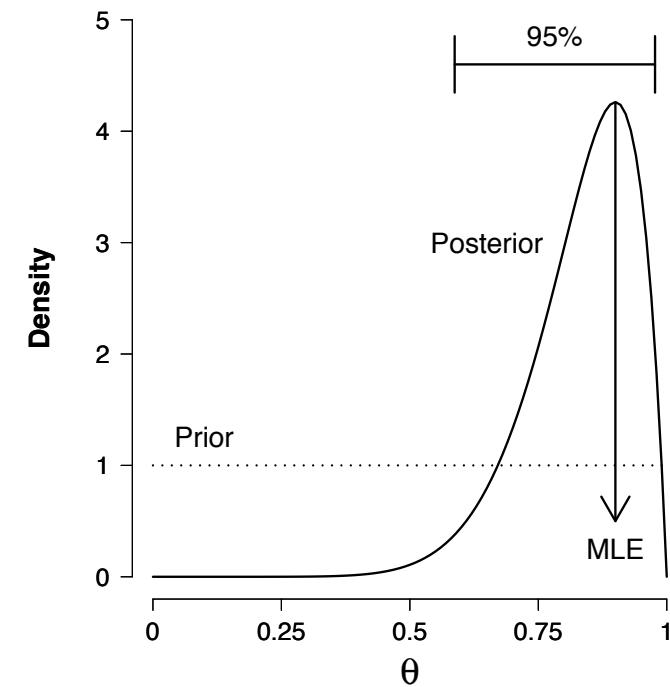
Relationship between MLE and Bayesian data analysis

An example from Lee & Wagenmakers (2012)

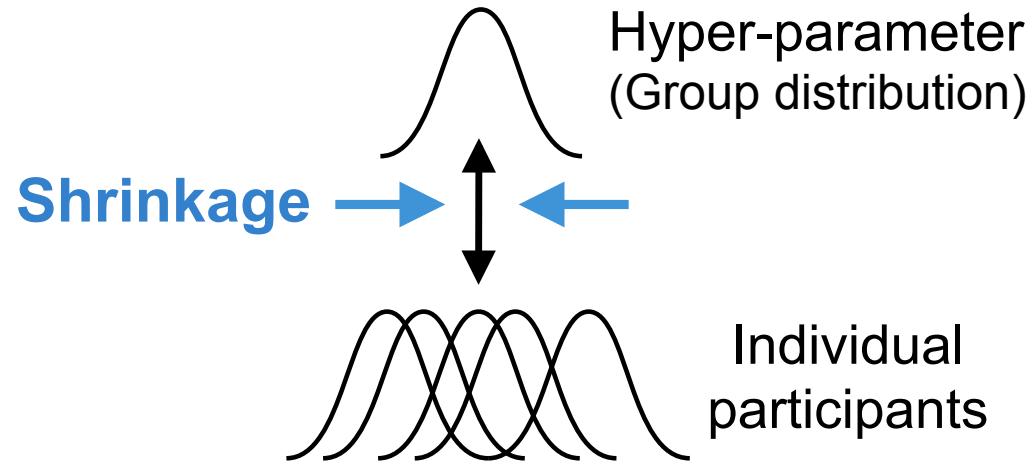
- *A test with 10 questions*
- *Parameter of interest: your ability ($0 < \theta < 1$)*
- *Data: 9 correct answers out of 10 ($k=9$, $n=10$)*
- *Likelihood function: Bernoulli*

Bayesian $p(\theta | D) = \frac{p(D | \theta) p(\theta)}{p(D)}$

MLE $\hat{\theta} = k/n = 0.9$



Hierarchical Bayesian analysis



- Similarities & differences across participants
- Small amount of data from each subject

$$P(\Theta, \Phi | D)$$

Posterior

$$\Theta_i = \{\alpha_i, \beta_i, \gamma_i, \dots\}$$

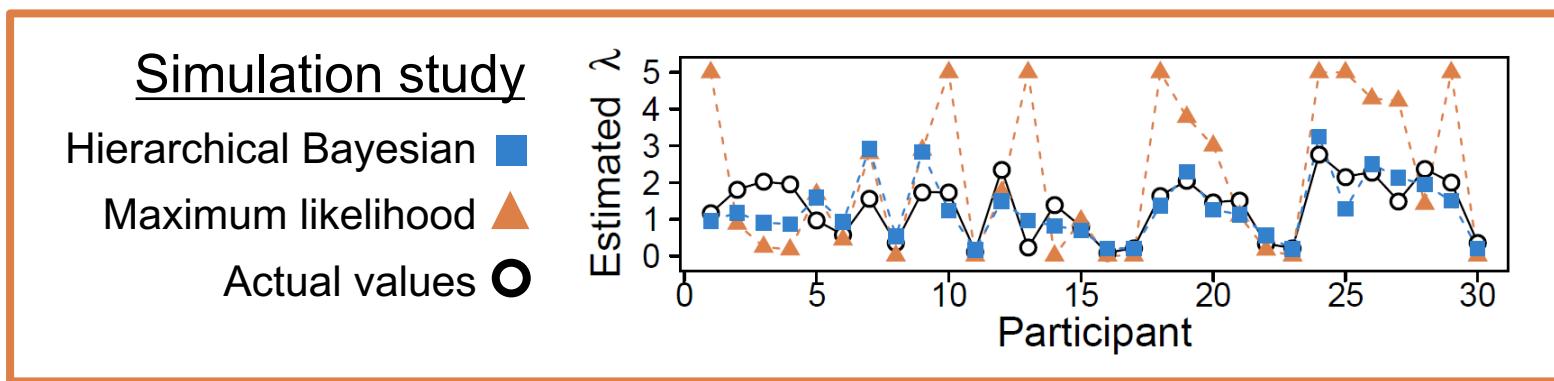
$$\Phi = \{\mu_\alpha, \mu_\beta, \mu_\gamma, \sigma_\alpha, \sigma_\beta, \sigma_\gamma, \dots\}$$

$$P(\Theta, \Phi | D) = \frac{P(D | \Theta, \Phi) P(\Theta, \Phi)}{P(D)} \propto P(D | \Theta) P(\Theta | \Phi) P(\Phi)$$

Likelihood Prior

Evidence

Hierarchical Bayesian data analysis leads to more accurate parameter estimates



Ahn et al. (2011, JNPE)

Limitations

- Assumption of a single hyper-group
- Shrinkage → regression toward the group mean

Ahn, W.-Y., Krawitz, A., Kim, W., Busemeyer, J. R., & Brown, J. W. (2011). A Model-Based fMRI Analysis with Hierarchical Bayesian Parameter Estimation. *Journal of Neuroscience, Psychology, and Economics*, 4(2), 95-110.

Okay sounds good... but how can we actually do Bayesian updating?

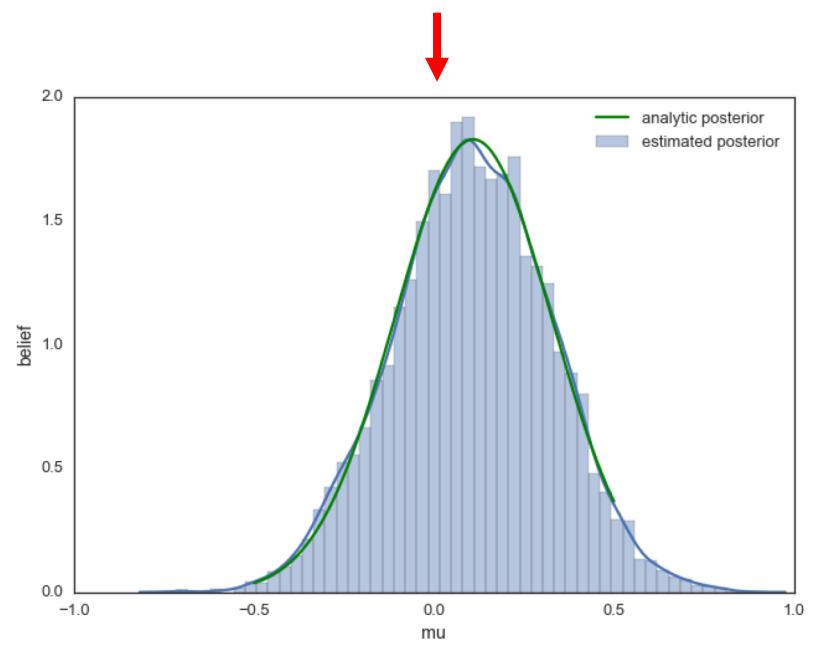
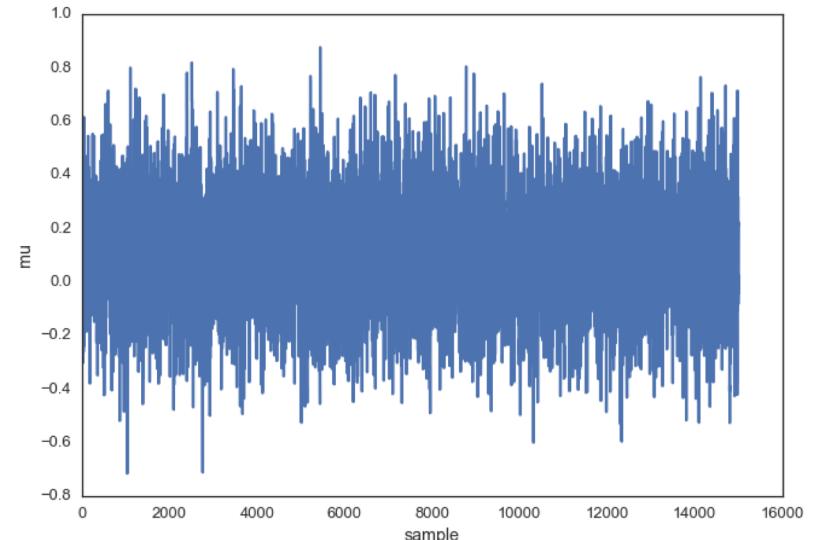
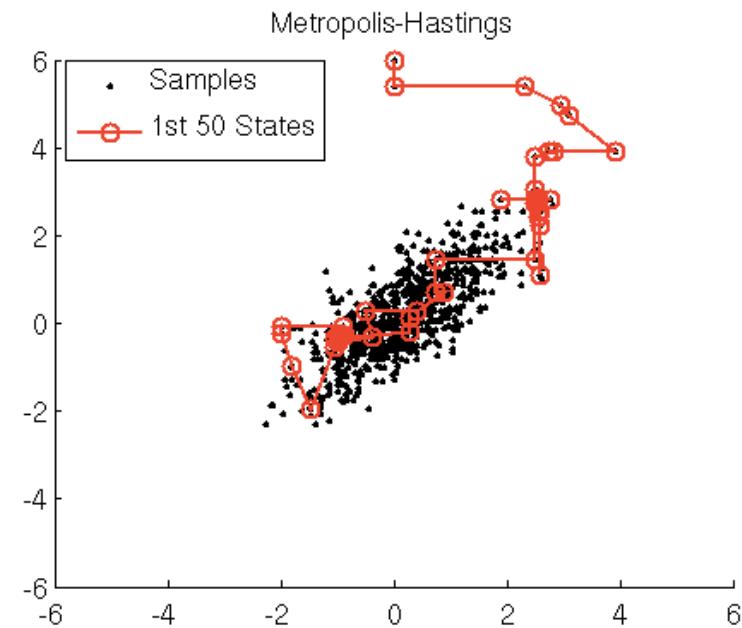
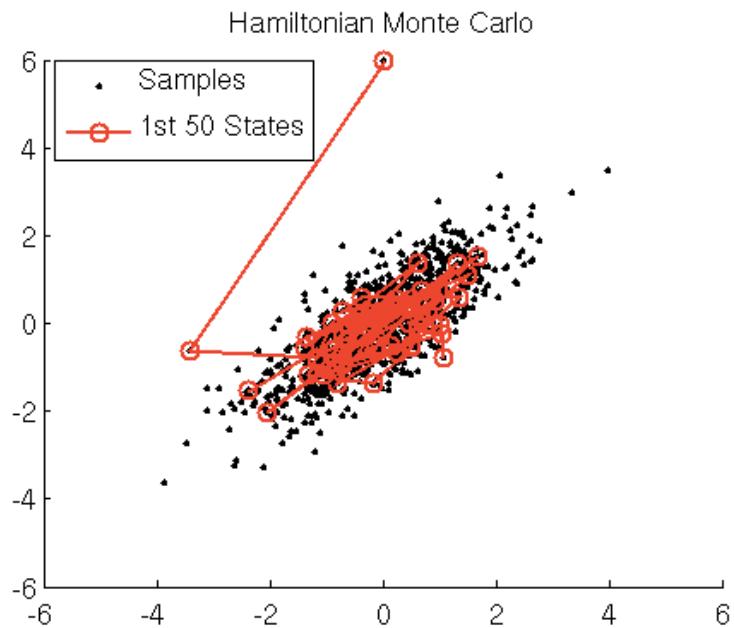
$$P(\Theta_i | D_i) = \frac{P(D_i | \Theta_i)P(\Theta_i)}{P(D_i)} = \frac{P(D_i | \Theta_i)P(\Theta_i)}{\int P(D_i | \Theta_{i'})P(\Theta_{i'})d\Theta_{i'}}$$

1. *Analytical solutions (often don't exist)*
2. *Grid approximation*
3. *Variational Bayes*
4. *Markov Chain Monte Carlo (MCMC)*

Markov Chain Monte Carlo (MCMC)

Approximate posterior distributions by drawing samples from the posterior distributions (after reaching an equilibrium).

Some samplers are more efficient than others!



<http://twiecki.github.io/blog/2015/11/10/mcmc-sampling/>

Tools for MCMC sampling

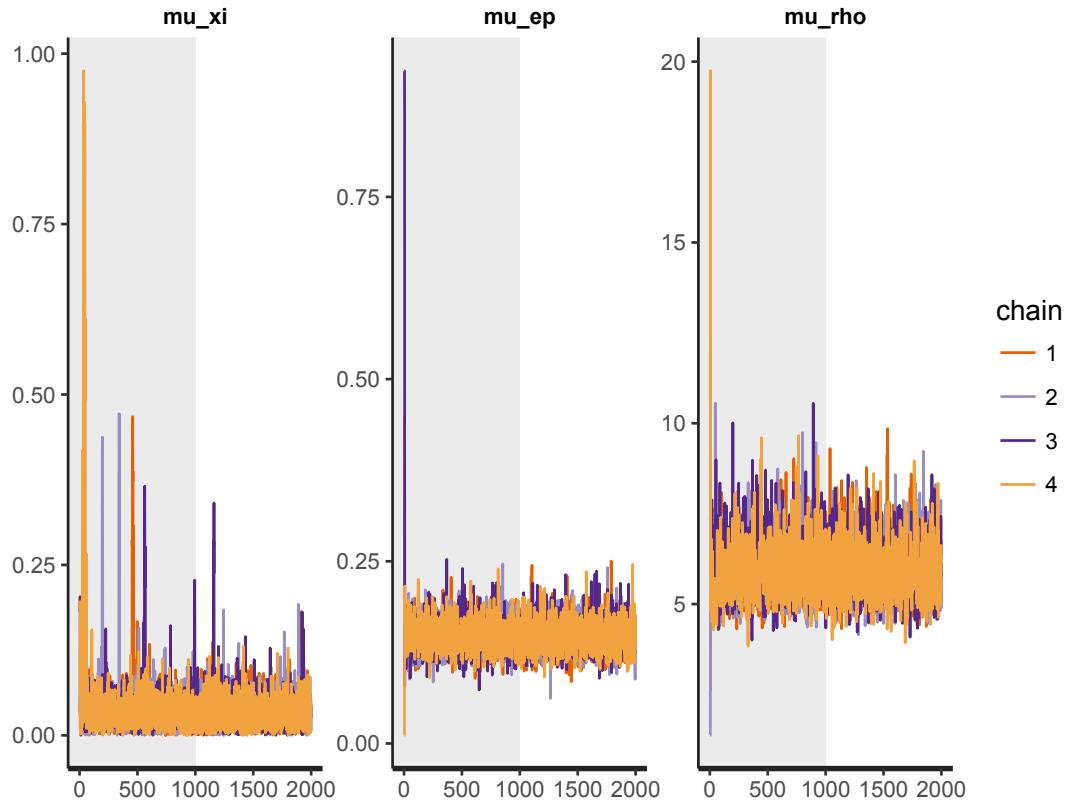
1. *Program your own!?*
2. *User-friendly packages*
 - *WinBUGS/OpenBUGS*
 - *JAGS*
 - *PyMC3 (in Python)*
 - *Stan*
3. *hBayesDM utilizes Stan (<http://mc-stan.org/>), which is well-maintained by a team led by Andrew Gelman.*
4. *Stan & hBayesDM support VB as well!*

Outline

- *Understand the concepts of computational modeling*
- *Know popular and cutting-edge methods for fitting computational models / Understand the concept of Bayesian data analysis*
- ***Things to know when performing MCMC sampling***
- *Use hBayesDM package to model several tasks*

Things to know when performing MCMC sampling

“Beware: *hBayesDM* can be dangerous!”



Check the convergence!

- Number of chains → at least 4
- Burn-in (warm-up) samples
- Convergence index (e.g., \hat{R}) → Close to 1.00
At least less than ~1.10

```
plot(output1, type="trace", inc_warmup=T)
```

Outline

- *Understand the concepts of computational modeling*
- *Know popular and cutting-edge methods for fitting computational models*
- *Understand the concept of Bayesian data analysis*
- *Evaluate outputs from Bayesian data analysis*
- ***Use hBayesDM package to model several tasks (Part II)***
 - *If you learn how to model one task, you know how to model all!*