

Workshop on the HGF-toolbox

Tore Erdmann, Lilian Weber

Zurich Computational Psychiatry Course
14.09.2018



Overview

- General framework
- Analysis workflow
- Intro to toolbox
- Response models
- Start with exercises

08.30-09.30

- Break


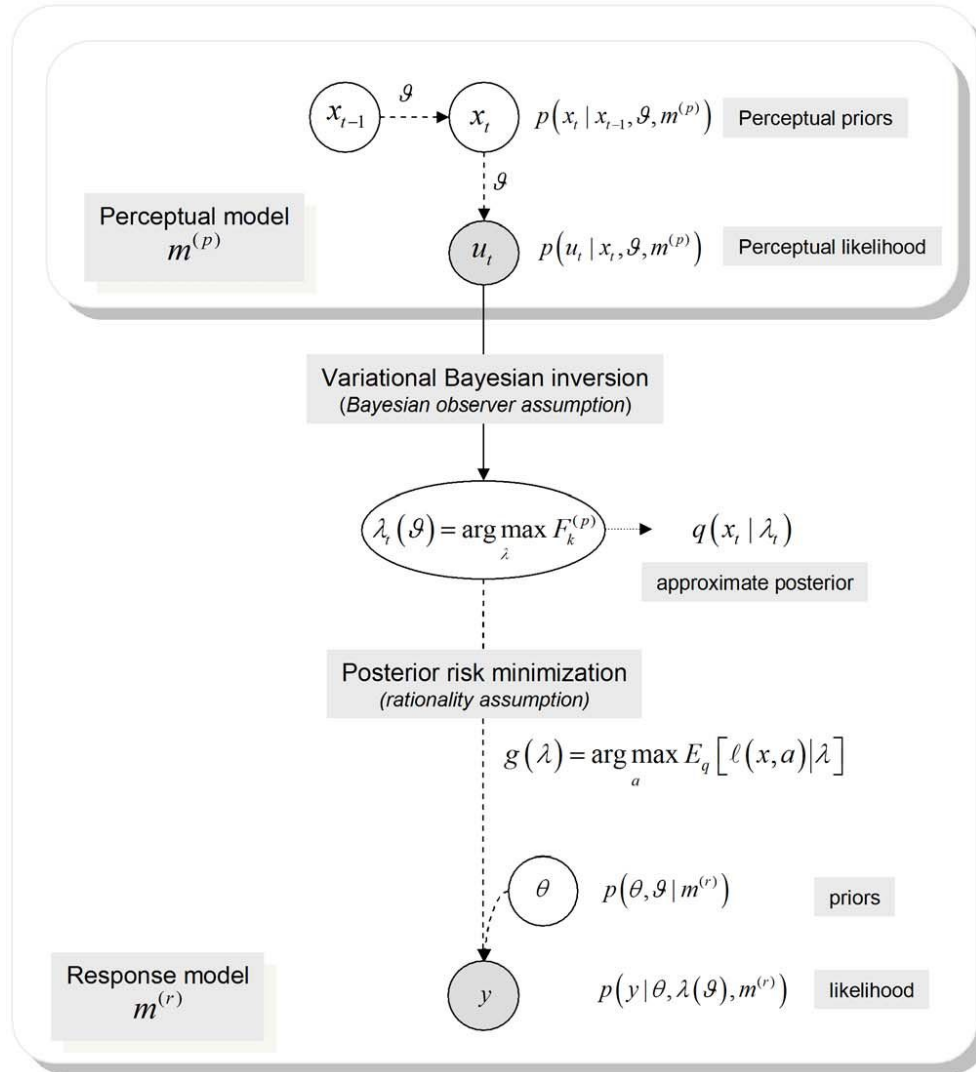
09.30-10.00

- Continue exercises
- Discussion / questions

10.00-11.30

Introduction

- We want to understand learning
- Learning = belief updating
- **Beliefs** here are (gaussian) probability distributions, and are updated based on inputs and a **perceptual model**
- **Responses** at each time are based on the current belief, as specified by a **response model**
- Together, these models provide a **generative model** of the responses of the learner, which **allows** us to do **inference on the model parameters**




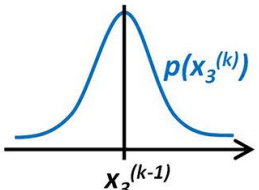
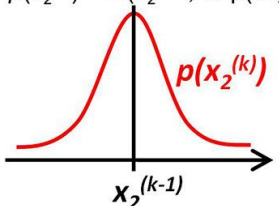
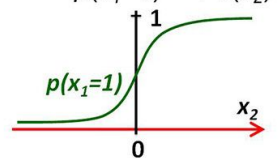
(Daunizeau et al., 2010)

Example: Reversal learning paradigm

- Two slot machines: Each trial, the participant makes a binary choice and is rewarded or not rewarded
- There is always one good machine, having 0.8 probability to give reward and can be reversed
- To do well, the agent needs to estimate (and update following a reversal) the probability for winning with each machine
- Tracking a single probability is enough because we have $P(\text{right is better}) = 1 - P(\text{left is better})$
- Initially the agent is uncertain about which is better, which can be expressed as the belief $P(\text{left is better}) = .5$

The generative model at the base of the learner's perceptual model



State of the world	Model
Log-volatility \mathbf{x}_3 of tendency	$p(x_3^{(k)}) \sim N(x_3^{(k-1)}, \vartheta)$ Gaussian random walk with constant step size ϑ 
Tendency \mathbf{x}_2 towards category "1"	$p(x_2^{(k)}) \sim N(x_2^{(k-1)}, \exp(\kappa x_3 + \omega))$ Gaussian random walk with step size $\exp(\kappa x_3 + \omega)$ 
Stimulus category \mathbf{x}_1 ("0" or "1")	$p(x_1=1) = s(x_2)$ $p(x_1=0) = 1-s(x_2)$ Sigmoid transformation of x_2 

Example: Reversal learning paradigm

At the bottom of the generative model, the learners observations are produced by:

$$u^{(k)} \sim \text{Ber}(x_1^{(k)})$$

which leads to these updates for the belief about the latent process,

$$\mu_2^{(k)} = \mu_2^{(k-1)} + \frac{1}{\pi_2^{(k)}} \delta_1^{(k)}$$

and which directly implies the degree of belief for the next observation:

$$\mu_1^{(k)} = s(\mu_2^{(k)})$$

Example: Reversal learning paradigm

This update is similar in form to a Rescorla-Wagner update: it adjusts the old prediction by a weighted prediction error:

$$\mu_2^{(k)} = \mu_2^{(k-1)} + \frac{1}{\pi_2^{(k)}} \delta_1^{(k)}$$

However, the weighting factor for the HGF is dynamic and reflects different kinds of uncertainty:

$$\frac{1}{\pi_2^{(k)}} = \frac{1}{\underbrace{\sigma^{(k-1)}}_{\text{estimation uncertainty}} + t^{(k)} \underbrace{\exp(\kappa \mu_3^{(k-1)} + \omega)}_{\text{estimated volatility of the environment (affected by parameters)}} + \underbrace{\frac{1}{\hat{\mu}_1^{(k)} (1 - \hat{\mu}_1^{(k)})}}_{\text{irreducible uncertainty about the outcome}}}$$

Update equations for expectations

<p>Level 3</p>	<div style="display: flex; justify-content: space-between; align-items: center;"> <div style="flex: 1;"> $\Delta\mu_3 = \sigma_3 \cdot \frac{\kappa}{2} \cdot w_2 \cdot \delta_2$ <div style="border: 1px solid black; padding: 5px; margin-top: 10px; width: fit-content;"> <p style="color: red;">Expectation update</p> <p style="color: green;">(Unweighted) learning rate</p> <p style="color: orange;">Weighting factor</p> <p style="color: blue;">Prediction error</p> </div> </div> <div style="flex: 1; text-align: right;"> <p>with</p> $\Delta\mu_3 = \mu_3^{(k)} - \mu_3^{(k-1)}$ $\sigma_3 = \sigma_3^{(k)}$ $w_2 = \frac{e^{\kappa\mu_3^{(k-1)} + \omega}}{\sigma_2^{(k-1)} + e^{\kappa\mu_3^{(k-1)} + \omega}}$ $\delta_2 = \frac{\sigma_2^{(k)} + \left(\mu_2^{(k)} - \mu_2^{(k-1)}\right)^2}{\sigma_2^{(k-1)} + e^{\kappa\mu_3^{(k-1)} + \omega}} - 1$ </div> </div>
<p>Level 2</p>	<div style="display: flex; justify-content: space-between; align-items: center;"> <div style="flex: 1;"> $\Delta\mu_2 = \sigma_2 \cdot \delta_1$ </div> <div style="flex: 1; text-align: right;"> <p>with</p> $\Delta\mu_2 = \mu_2^{(k)} - \mu_2^{(k-1)}$ $\sigma_2 = \sigma_2^{(k)}$ $\delta_1 = \mu_1^{(k)} - s\left(\mu_2^{(k-1)}\right)$ </div> </div>

Update equation for posterior means

$$\Delta\mu_i^{(k)} = \frac{1}{2} \kappa_{i-1} v_{i-1}^{(k)} \frac{\hat{\pi}_{i-1}^{(k)}}{\pi_i^{(k)}} \delta_{i-1}^{(k)}$$

$$\Rightarrow \boxed{\Delta\mu_i \propto \frac{\hat{\pi}_{i-1}}{\pi_i} \delta_{i-1}}$$

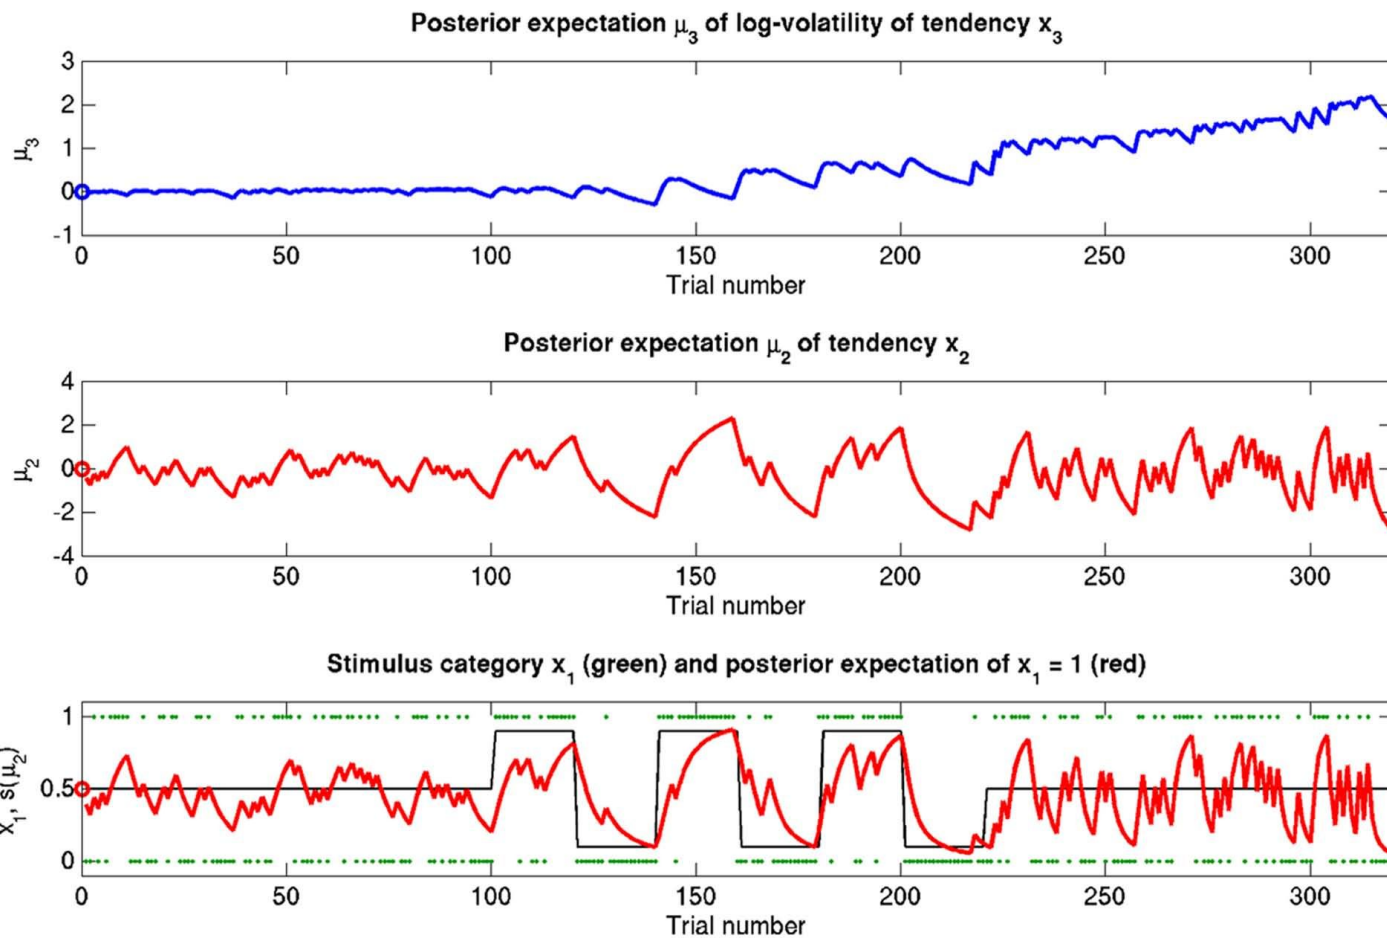
with

$\hat{\pi}_{i-1}$	Precision of the prediction onto the level below
-------------------	--




π_i	Posterior precision at the current level
---------	--

δ_{i-1}	Prediction error for the volatility of the level below
----------------	--


Binary HGF



Binary HGF: reduced theta



Binary HGF: reduced omega



Adding a response model

- Ok, so we have a model about how to update beliefs
- Now we need another model to map beliefs to responses
- Some simple suggestions for binary:
 - If $p > .5$ do action 1
 - Sample action $\sim \text{bernoulli}(p)$
 - Binary softmax with ‘decision temperature’ \rightarrow allows modeling exploration / noisiness of responses
- Response models can also be regression models, drift diffusion, LBAs ...

Types of input- and response variables

- Think about the data analysis before starting the data collection
- Binary inputs/responses is not as informative as graded inputs/responses

	Input: binary	Input: continuous
Output: binary	Input: blue/green card Choice: yes/no	Input: pitch Choice: yes/no
Output: continuous	Input: blue/green card Choice: log(RT)	Input: pitch Choice: log(RT)

Analysis workflow

1. Specify model (and design experiment)
2. Setting priors / prior predictive checking
3. (Get data)
4. Fit model
5. Check model
6. Good enough? Publish. Else, adjust model and go to 1.

HGF-toolbox

- Functions for fitting, simulation, diagnostics
- Various choices for perceptual and response models (and can be extended):

Belief trajectory:	Response models:
Binary	Binary (square sigmoid, binary softmax)
Categorical	Categorical (softmax)
Continuous	Continuous (gaussian, linear regression)

Important functions

- Fitting the model

```
fit = tapas_fitModel(responses, inputs, <prc_model_config>,  
                    <obs_model_config>);
```

- Simulation from the model

```
sim = tapas_simModel(inputs, <prc_model>, prc_model_parameters, ...  
                    <obs_model>, obs_model_parameters);
```

- Plotting belief trajectories / simulated responses

```
tapas_hgf_plotTraj(fit);
```

```
tapas_hgf_binary_plotTraj(fit);
```

Using `tapas_fitModel`

This call estimates parameters for given inputs and responses:

```
fit = tapas_fitModel(responses, inputs, ...  
                    <prc_model_config>, ...  
                    <obs_model_config>, ...  
                    <opt_algo>);
```

Annotations for configuration strings:

- For `<prc_model_config>`:
 - 'tapas_hgf'
 - 'tapas_hgf_binary'
 - 'tapas_rw_binary'
- For `<obs_model_config>`:
 - 'tapas_hgf_categorical'
- For `<opt_algo>`:
 - 'tapas_gaussian_obs_config'
 - 'tapas_softmax_binary_config'

Using `tapas_simModel`

This call simulates responses for given inputs and parameters:

```
sim = tapas_simModel(inputs, ...  
                    <prc_model>, ...  
                    prc_parameters, ...  
                    <obs_model>, obs_parameters);
```

Diagram illustrating the structure of the `tapas_simModel` function call:

- `<prc_model>` is associated with the list: `'tapas_hgf'`, `'tapas_hgf_binary'`, `'tapas_rw_binary'`, `'tapas_hgf_categorical'`.
- `<obs_model>` is associated with the list: `'tapas_gaussian_obs'`, `'tapas_softmax_binary'`.

Model comparison

- When to choose between models?
 - Finding the right model
 - Justifying your model
- How to choose a model?
 - Do prior predictive checks
 - Only keep sensible models, you don't need try every possible model
 - Do posterior predictive checks: is the model missing something?
 - Compare models using the model evidence $p(y|M)$
- ``tapas_fitModel`` computes the model evidence for you
 - `est.optim.LME`
 - Fixed-effects analysis: for each model sum the LME values of each participant and take exp of the difference
 - Random-effects analysis: use ``spm_BMS`` from the SPM toolbox

On to the exercises!